

# 1. Project Overview

**Objective:** Segment retail customers into meaningful groups using transactional data, enabling targeted marketing, retention strategies, and revenue optimization.

**Dataset:** Online Retail Transaction Data (Invoice-level)

**Approach:**

- Rigorous data cleaning
- Customer-level feature engineering (RFM + behavioral)
- Proper scaling and validation
- Unsupervised learning (KMeans)
- Statistical and business interpretation\*

## 2. Imports & Configuration

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture
from sklearn.metrics import silhouette_score, adjusted_rand_score

import hdbscan

RANDOM_STATE = 53
np.random.seed(RANDOM_STATE)
```

## 3. Data Loading

```
data = pd.read_csv('customer_segmentation.csv', encoding='latin-1')
data['InvoiceDate'] = pd.to_datetime(data['InvoiceDate'])
```

## 4. Data Cleaning

```
# Remove cancelled invoices
data = data[~data["InvoiceNo"].astype(str).str.startswith("C")]

# Remove missing customers
data = data.dropna(subset=["CustomerID"])

# Remove invalid quantities and prices
data = data[(data["Quantity"] > 0) & (data["UnitPrice"] > 0)]

# Fix product descriptions (StockCode → mode Description)
desc_map = (
    data.dropna(subset=["Description"])
    .groupby("StockCode")["Description"]
    .agg(lambda x: x.value_counts().index[0])
)

data["Description"] = data["Description"].fillna(
    data["StockCode"].map(desc_map)
)

# Remove non-product codes
noise_codes = ["POST", "D", "M", "DOT", "BANK CHARGES"]
data = data[~data["StockCode"].isin(noise_codes)]

# Monetary value
data["TotalPrice"] = data["Quantity"] * data["UnitPrice"]
```

## 5. Feature Engineering

### 5.1 RFM Features

```
snapshot_date = data["InvoiceDate"].max() + pd.Timedelta(days=1)

rfm = data.groupby("CustomerID").agg({
    "InvoiceDate": lambda x: (snapshot_date - x.max()).days,
    "InvoiceNo": "nunique",
    "TotalPrice": "sum"
}).reset_index()

rfm.columns = ["CustomerID", "Recency", "Frequency", "Monetary"]
```

### 5.2 Behavioral Features

```
behavior = data.groupby("CustomerID").agg({
    "Quantity": "sum",
    "StockCode": "nunique"
```

```
}).reset_index()

behavior.columns = ["CustomerID", "TotalQuantity", "UniqueProducts"]

customer_df = rfm.merge(behavior, on="CustomerID")
```

## 6. Feature Scaling

```
features = [
    "Recency",
    "Frequency",
    "Monetary",
    "TotalQuantity",
    "UniqueProducts"
]

scaler = StandardScaler()
X_scaled = scaler.fit_transform(customer_df[features])
```

## 7. Baseline Segmentation


### 7.1 Model Selection

```
inertia = []
sil_scores = []

for k in range(2, 11):
    km = KMeans(n_clusters=k, random_state=RANDOM_STATE)
    labels = km.fit_predict(X_scaled)
    inertia.append(km.inertia_)
    sil_scores.append(silhouette_score(X_scaled, labels))

plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(range(2,11), inertia, marker="o")
plt.title("Elbow Method")

plt.subplot(1,2,2)
plt.plot(range(2,11), sil_scores, marker="o")
plt.title("Silhouette Scores")
plt.show()
```



## 7.2 Final K-Means

```
kmeans = KMeans(n_clusters=4, random_state=RANDOM_STATE)
customer_df["KMeans_Cluster"] = kmeans.fit_predict(X_scaled)

silhouette_score(X_scaled, customer_df["KMeans_Cluster"])

0.512926083590546
```

# 8 . Cluster Stability Testing

## 8.1 Bootstrap Stability

```
n_bootstraps = 30
ari_scores = []

base_labels = customer_df["KMeans_Cluster"].values

for i in range(n_bootstraps):
    idx = np.random.choice(len(X_scaled), len(X_scaled), replace=True)
    X_sample = X_scaled[idx]

    km = KMeans(n_clusters=4, random_state=i)
    labels = km.fit_predict(X_sample)

    ari = adjusted_rand_score(base_labels[idx], labels)
    ari_scores.append(ari)

print(f"Mean ARI: {np.mean(ari_scores):.3f}")
print(f"Std ARI : {np.std(ari_scores):.3f}")

Mean ARI: 0.737
Std ARI : 0.305
```

## Interpretation rule (state this in report):

- ARI > 0.75 → highly stable
- 0.5–0.75 → acceptable
- < 0.5 → unreliable

## 9. Alternative Model I: Gaussian Mixture Model (GMM)

```
gmm = GaussianMixture(  
    n_components=4,  
    covariance_type="full",  
    random_state=RANDOM_STATE  
)  
  
customer_df["GMM_Cluster"] = gmm.fit_predict(X_scaled)  
  
# Soft membership confidence  
gmm_probs = gmm.predict_proba(X_scaled)  
customer_df["GMM_MaxProb"] = gmm_probs.max(axis=1)  
  
silhouette_score(X_scaled, customer_df["GMM_Cluster"])  
0.107471585490664
```

## 10. Alternative Model II: HDBSCAN (Density-Based)

```
hdb = hdbscan.HDBSCAN(  
    min_cluster_size=30,  
    min_samples=15  
)  
  
customer_df["HDBSCAN_Cluster"] = hdb.fit_predict(X_scaled)  
  
customer_df["HDBSCAN_Cluster"].value_counts()  
  
{"columns":  
 [{"name": "HDBSCAN_Cluster", "rawType": "int64", "type": "integer"},  
 {"name": "count", "rawType": "int64", "type": "integer"}], "ref": "dbcfabbb-  
 b015-48f8-9df3-c84661b7f922", "rows": [[[-1, "2818"], [4, "568"],  
 ["5", "414"], [2, "328"], [0, "85"], [3, "82"], [1, "39"]], "shape":  
 {"columns": 1, "rows": 7}}}
```

- $-1 \rightarrow$  noise / unsegmentable customers
- Indicates niche or anomalous buyers

## 11. Comparative Analysis

```
comparison = customer_df.groupby("KMeans_Cluster")[features].mean()
comparison

{
  "columns": [
    {"name": "KMeans_Cluster", "rawType": "int32", "type": "integer"},  

    {"name": "Recency", "rawType": "float64", "type": "float"},  

    {"name": "Frequency", "rawType": "float64", "type": "float"},  

    {"name": "Monetary", "rawType": "float64", "type": "float"},  

    {"name": "TotalQuantity", "rawType": "float64", "type": "float"},  

    {"name": "UniqueProducts", "rawType": "float64", "type": "float"}], "ref": "1  

0e7c72c-b44c-45bb-f08621f1039a", "rows":  

[[["0", "46.32494684620836", "3.2955350815024804", "1138.5100715804394", "7  

01.785258681786", "49.324238128986536"],  

["1", "16.29705215419501", "14.22222222222221", "6975.729750566893", "404  

9.278911564626", "207.81179138321994"],  

["2", "25.764705882352942", "69.88235294117646", "111849.4594117647", "644  

81.64705882353", "592.7058823529412"],  

["3", "249.92504743833018", "1.560721062618596", "547.5707789373814", "280  

.0123339658444", "24.040796963946867]]}, "shape": {"columns": 5, "rows": 4}  

}  

pd.crosstab(  

    customer_df["KMeans_Cluster"],  

    customer_df["GMM_Cluster"],  

    normalize="index"
)  

{
  "columns": [
    {"name": "KMeans_Cluster", "rawType": "int32", "type": "integer"},  

    {"name": "0", "rawType": "float64", "type": "float"},  

    {"name": "1", "rawType": "float64", "type": "float"},  

    {"name": "2", "rawType": "float64", "type": "float"},  

    {"name": "3", "rawType": "float64", "type": "float"}], "ref": "98dbf94c-996d-  

4f18-badd-818832f80dfa", "rows":  

[[["0", "0.5262225372076541", "0.1708008504606662", "0.011339475549255847"  

, "0.2916371367824238"],  

["1", "0.0022675736961451248", "0.8480725623582767", "0.14965986394557823  


, "0.0"], ["2", "0.0", "0.0", "1.0", "0.0"],  

["3", "0.28178368121442127", "0.0", "0.016129032258064516", "0.70208728652  

75142]], "shape": {"columns": 4, "rows": 4}}
}
```

## 12. Visualization

```
plt.figure(figsize=(8,6))
sns.scatterplot(
    data=customer_df,
    x="Monetary",
    y="Frequency",
    hue="KMeans_Cluster",
    palette="viridis",
    alpha=0.7
)
plt.title("Customer Segments (KMeans)")
plt.show()
```



## 13. Business Interpretation

Cluster	Description
0	Bulk low-frequency buyers

Cluster	Description
1	Recent low spenders
2	High-value loyal customers
3	At-risk customers

## 14. Production Pipeline Design

### Inference Function

```
FINAL_MODEL = kmeans    # baseline, stable, interpretable
FINAL_SCALER = scaler
FINAL_FEATURES = features

def assign_customer_segment(df):
    X = FINAL_SCALER.transform(df[FINAL_FEATURES])
    return FINAL_MODEL.predict(X)
```

### Artifacts to Persist

- scaler.pkl
- kmeans\_model.pkl
- Feature schema JSON

```
import joblib
import json
from pathlib import Path

ARTIFACT_DIR = Path("../artifacts")
ARTIFACT_DIR.mkdir(exist_ok=True)

joblib.dump(FINAL_SCALER, ARTIFACT_DIR / "scaler.pkl")
joblib.dump(FINAL_MODEL, ARTIFACT_DIR / "kmeans.pkl")
joblib.dump(gmm, ARTIFACT_DIR / "gmm.pkl")

['..\\"artifacts\\gmm.pkl']

feature_schema = {
    "features": FINAL_FEATURES,
    "description": {
        "Recency": "Days since last purchase",
        "Frequency": "Number of invoices",
        "Monetary": "Total spend",
        "TotalQuantity": "Total items purchased",
        "UniqueProducts": "Distinct products purchased"
    }
}
```

```
with open(ARTIFACT_DIR / "feature_schema.json", "w") as f:  
    json.dump(feature_schema, f, indent=4)
```